Sandboxing is (the) shit!

Jonathan Brossard (Toucan System)



Who am I?

- Security researcher, publishing since 2005.
- Past research: vulnerabilities in BIOSes, Microsoft Bitlocker, Truecrypt, McAfee Endpoint (Defcon 2008), PMCMA debugger (Blackhat USA 2011), « Rakshasa » supply chain backdoor PoC (Blackhat 2012), 2 SAP notes (2013).
- Speaker/trainer at HITB, CCC, Ruxcon...
- Co-founder of the Hackito Ergo Sum and NoSuchCon research conferences (France).

Disclaimer: contains research

This was supposed to be a short research on finding/exploiting a few cool low level bugs in sandboxes.

It ended up leading to more questions than answers on my understanding of what the industry is doing in the AV/sandbox space.

If you have better understanding, I'd really like if you took the time to explain me

(endrazine@gmail.com,+PGP).



Disclaimer (rephrased)

WTF is the AV industry doing? Well, I'm not so sure I understand anymore ...

What's hot in the AV industry in 2013?



AV industry: 2013 trends

- Desktop AV is essentially a thing of the past
- Focus moves technologies hopefully able to « detect 0days »[1] like <u>sandboxing</u>.

=> The new cool thing is emulation and sandboxing.

[1] Don't laugh yet.

How it all started... (/story telling)



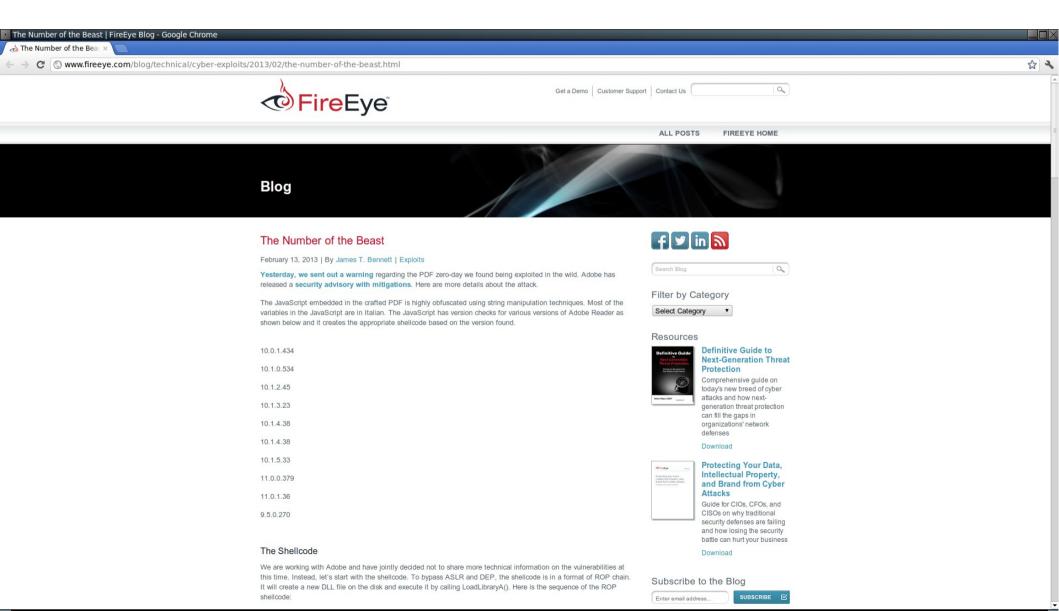
CVE-2013-0640 (Adobe Sandbox bypass)



Adobe Reader and Acrobat 9.x before 9.5.4, 10.x before 10.1.6, and 11.x before 11.0.02 allow remote attackers to execute arbitrary code or cause a denial of service (memory corruption) via a crafted PDF document, as exploited in the wild in February 2013.



CVE-2013-0640 (Adobe Sandbox bypass)



Their « analysis »:

Here is the sequence of the ROP shellcode:

msvcr100!fsopen()

msvcr100!write()

mvvcr100!fclose()

kernel32!LoadLibraryA()

kernel32!Sleep()

Upon loading the malicious library, it will enter a long sleep and ensure that the thread has not crashed because the whole stack in the thread is already manipulated for creating a ROP chain.



Their « analysis »:

Here is the sequence of the ROP shellcode:

msvcr100!fsopen()

msvcr100!write()

mvvcr100!fclose()

kernel32!LoadLibraryA()

kernel32!Sleep()



Upon loading the malicious library, it will enter a long sleep and ensure that the thread has not crashed because the whole stack in the thread is already manipulated for creating a ROP chain.

Their « analysis »:

Here is the sequence of the ROP shellcode:

msvcr100!fsopen()

msvcr100!write()

mvvcr100!fclose()

kernel32!LoadLibraryA()

kernel32!Sleep()



Upon loading the malicious library, it will enter a long sleep and ensure that the thread has not crashed because the whole stack in the thread is already manipulated for creating a ROP chain.

=> In trivial english, this is called bullshitting. They clearly have no idea what the exploit is trying to do here.

What I believe really happens in this case (wild guess)

Sleep 5 minutes to attempt bypass sanboxing detection:)

After all, it's a hardened exploit, found in the wild and the first of its kind to bypass Adobe sandboxing technology...

Limits of such technologies (imho)

- Good at finding artefacts (it's still « something »).
- Pretty bad at understanding what is actually happening inside the exploit.

That being said...





About

Being able to understand the way malware operate is the key to properly **fight** them. Cuckoo Sandbox helps you achieving this goal in an easy and automated fashion.

Read more »

Download

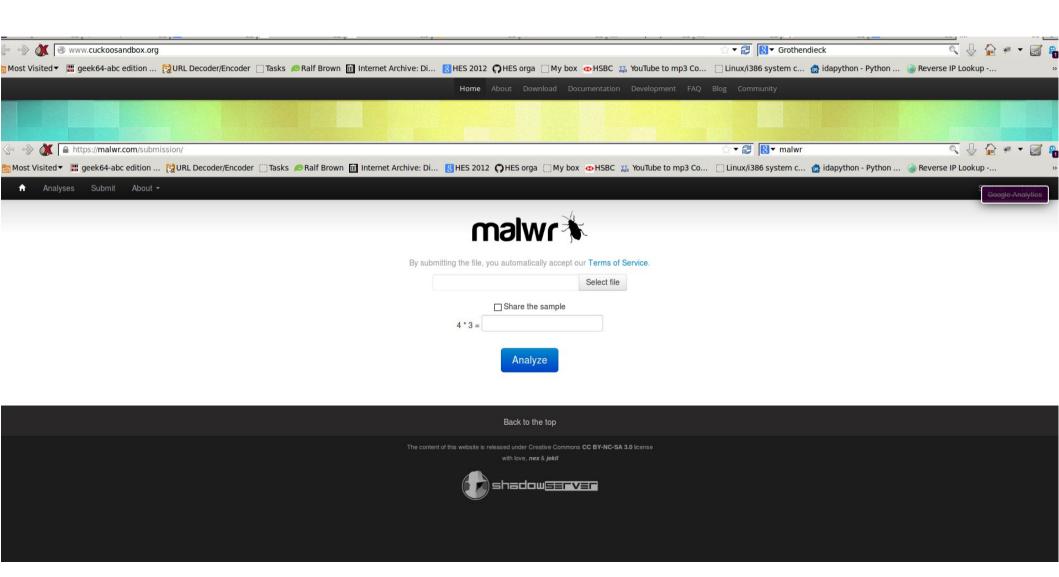
Cuckoo Sandbox is a completely **open source** solution, meaning that you can look at its internals, modify it and customize it at your will. Go on and download it to start tackling malware.

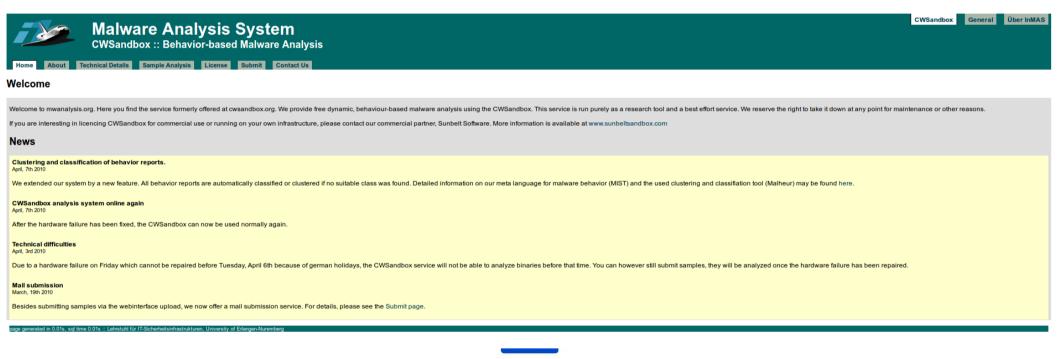
Read more »

Participate

Cuckoo Sandbox is a **community** effort. The only reason of it's growth and popularity is the people using it and contributing to it. Get in touch with the developers and with the users now!

Read more »





Back to the top

The content of this website is released under Creative Commons CC BY-NC-SA 3.0 license with love. nex & iekil





Anubis: Analyzing Unknown Binaries

Home Advanced Clustering News About Sample Reports



Welcome to Anubis

Anubis is a service for analyzing malware.

Submit your Windows executable or Android APK and receive an analysis report telling you what it does. Alternatively, submit a suspicious URL and receive a report that shows you all the activities of the Internet Explorer process when visiting this URL.

Ewitter. Want notifications about Anubis downtimes and/or updates? Follow us on twitter.

-Announcement



We are proud to present our most recent substantial extension to Anubis: the analysis of Android APKs (codename Andrubis)!

Like the core-Anubis does for Windows PE executables. Andrubis executes Android apps in a sandbox and provides a detailed report on their behavior, including file access, network access, crypto operations, dynamic code loading and information leaks. In addition to the dynamic analysis in the sandbox, Andrubis also performs static analysis, yielding information on e.g. the app's activities, services, required external libraries and actually required permissions.

To analyze apps straight away from your smartphone, check out our experimental submission app! Available in the Play Store soon.

09.10.2012 We are currently migrating to new hardware. Please report any service problems you experience!

30.05.2012 You can now also submit Android APKs!

16.02.2012 Five years Anubis!

05.07.2010 We have improved our analysis of network dumps. Extended DNS data (such as multiple DNS replies) are now available in the analysis reports.

02.07.2010 Dionaea/Nepenthes can again automatically upload samples to Anubis. We will reply with an analysis report!

01.06.2010 The DII-analysis has been improved. Simply upload a dynamically linked library file for Windows, and we'll try to figure out how to analyze it best!

01.03.2010 We have vastly improved analysis performance of the sandbox. You should now get more analysis results for the same execution duration!

Choose the subject for analysis

For analyzing Javascript and Flash files try Wepawet.

• File: (max. 8MB)

O URL:

Choose the file that you want to analyze. The file must be a Windows executable or Android APK. (details)

Browse... No file selected.

Choose the URL that you want to analyze. The URL will be analyzed in Internet Explorer.

Note: We will not analyze a binary that you provide via this URL. We will merely use a browser to check the given URL for a possible drive-by download or similar attack!

Note to self: I don't find quite reasonable to add to your corporate network something nobody really understands.



Dear @FireEye, is there a chance you will let independant researchers look at your appliance? Lack of transparency=not good:) /cc @taviso







@endrazine @HaifeiLi @taviso surprise plays a part in security, if the adversary cannot examine your defences before encountering them...





@HaifeiLi @endrazine @taviso if you want to remain effective as a security vendor, having working tech basically requires keeping it secret.



The whole concept of sandboxing vendors is to not have the perceived enemy take a look at the technology. Ok, agreed.

It also means no third party assessment has been done by the security community...

In real life, having software due dilligence done by the community has proved to be a good thing for the quality of the said software.

See similar requests from Tavis Ormandi and Pipacs to have a look at Bromium's technology...

Note: well, they're not Bromium clients, so we have a problem... as an industry, really.

Note 2 : Afaik, Bromium has researchers like Nergal and Jarred Demott. Who of this caliber works for FireEye really?





A petition to get a Bromium demo for @taviso and Pipacs. Pipacs has agreed to use a 10 year old copy of IDA to handicap his breaking it.



8:40 PM - 6 Aug 13

Reply to @grsecurity @taviso



Kostya Kortchinsky @crypt0ad @grsecurity @taviso can I play too??

Details

6 Aug

Room for problems (research leads)



Room for problems part I: general design/architecture



What's the trend, perceived objective like

Current « genious » idea :

- Correlate/share more data to create information asymetry.
- To do that, most (all I've seen allow it, at least in non default mode) solutions now <u>allow a malware to connect back to the internet[*]</u>.
- [*] Idea being to correlate DNS/binary checksum informations over « campaigns » of attacks in the time.

A few facts on this...

- The whole corporate strategy over the past 15 years has been to segregate LANS, DMZs and the internet.
- Now you give a temporary shell to the attacker at network perimeter (proxies, mail gateways, wherever such sandboxing solutions exist)...

- ... and what happens to your DNS? To your http proxy cache?

The Katsuni-Kaminsky attack (having it both ways)

- Attacker can run a malware inside a sandbox.
- Sandbox allows attacker to connect back to the internet.
- Corporate DNS server is used as a recursive DNS server.
- Attacker has it both ways and can synchronize arbitrary spoofed packets emission from both inside and outside the network.

=> That's gonna be very « safe » for sure...

Not to mention...

- What happens if the malware, from inside the sandbox manages to attack other networks (say crowdstrike!)?
- What happens if the malware can send back a modified copy of itself to the same network (smtp?) for more analysis, and more sandbox cpu time?

=> It's all about implementation details really.

Is this « wormable » (yet)?

- FireEye claims to work with 30 % of the TOP 100 Companies. That makes it not so hard to find...
- Their strategy is to synchronize malware information sharing... (allow exploits to drop exes on the sandboxes...?!?)
- They cover you « 360 », from mail gateways to http proxy file downloads, etc.
- Ok, so how exactly do you prevent one malware to get endless free execution time inside your different sandboxes around the world?

Room for problems part II: Turning lame bugs into sandbox Oracles



The problem

- Many online malware scanning engines run qemu (+ some various instrumentation and automation custom software)
- User doesn't get to see anything from the scanning process.
- Now, what if an attacker has a qemu lame DoS or endless loop PoC ?

Turning lame bugs into sandbox Oracles

(aka : hacking online malware analysis tools... hrm)

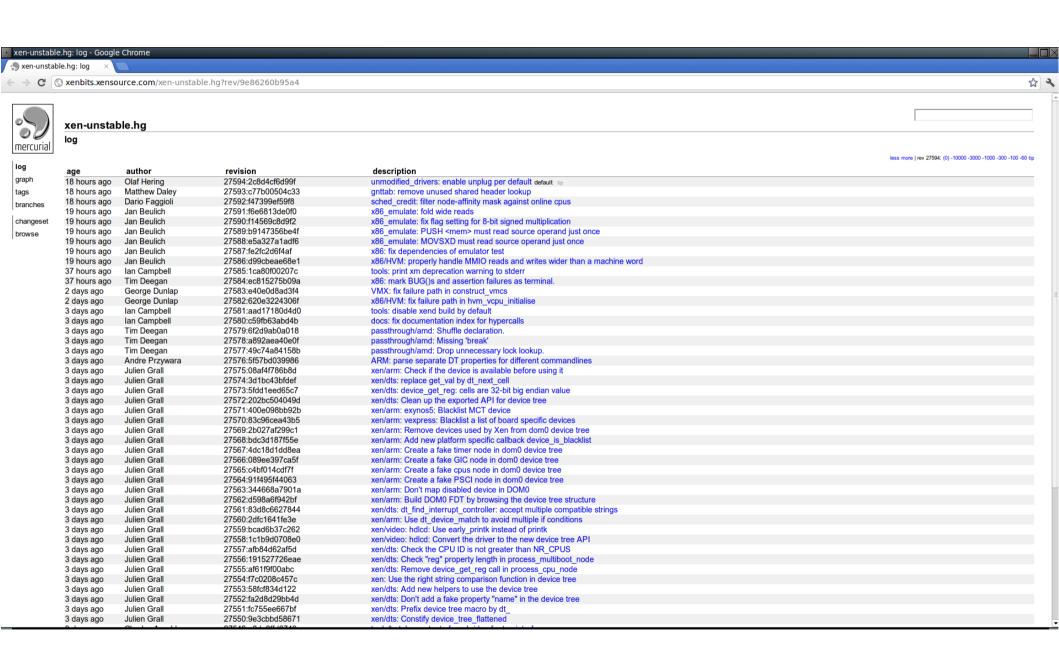
Room for problems part III: Such bugs do happen...



What degrees do projects like Xen or qemu really have in terms of security?

A fair question is the relative maturity of such technology, not when it comes to support legacy Oses or current Oses, but hostile malware trying to hurt them.

Exempli gratia: typical bugs reported (complexity, software maturity: format strings/symlinks or complex overflows?)



x86_emulate: MOVSXD must read source operand just once (Xen Unstable, 21/09/2013)

```
if ( (rc != X86EMUL OKAY) ||
1.109
                (regs.eax != 0x0000ffff) ||
1.110
                ((regs.eflags&0x240) != 0x200) ||
1.111
                (regs.eip != (unsigned long)&instr[2]) )
1.112 -
                (regs.eip != (unsigned long)&instr[2 + (ctxt.addr size
1.113 +
              goto fail;
1.114
          printf("okay\n");
1.115
1.116
1.117 @@ -821,6 +877,8 @@ int main(int argc, char **argv)
              if (j == 2) break;
1.118
              memcpy(res, blowfish32_code, sizeof(blowfish32_code));
1.119
```

Ok, so you're at miscomputing EIP on basic instructions such as MOVSXD on 32b...

Note: this is totally exploitable imho btw.

How does that rank compared to real cpu bugs?

Intel® Pentium® Processor

Invalid Instruction Frratum Overview

Invalid Instruction Errata Home

Software Backgrounder

Updated Nov. 20 1997

Software Vendor Statements

Erratum Technical Description

Updated Nov. 20 1997

Intel Identifies Workaround for the "Invalid Operand with Locked Compare Exchange 8Byte (CMPXCHG8B) Instruction" Erratum

Erratum Overview

On Friday, November 7th 1997, a number of reports were posted to the Internet implying the possibility of a new erratum on the Pentium® processors and Pentium® processors with MMX™ technology. An erratum is a design defect or error which may cause a product to deviate from published specifications. Based on the Internet reports our engineering team quickly jumped on this issue. Once we were able to reproduce the behavior we confirmed that an erratum does exist which is now named the "Invalid Operand with Locked CMPXCHG8B instruction" erratum. We were also able to identify the following:

That was in 1997... The instruction is far less used (CMPXCHG8B)

Qemu internals



Qemu goals

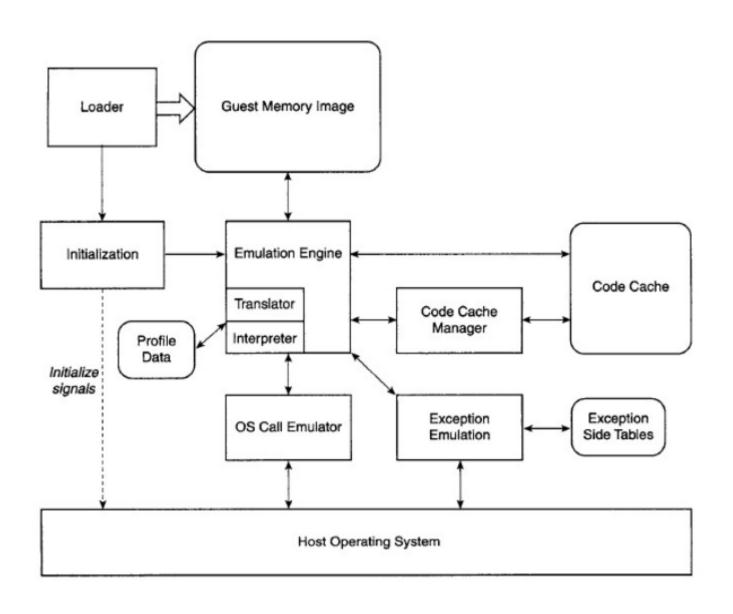
- (Fast) binary translation
- Binary code is translated into an IR, which is then executed by a virtual cpu, independently of host OS/arch.
- Super generic, super fast, super portable, super cool. Truely impressive.
- Super secure for the purpose of malware analysis though?

Qemu supports two modes

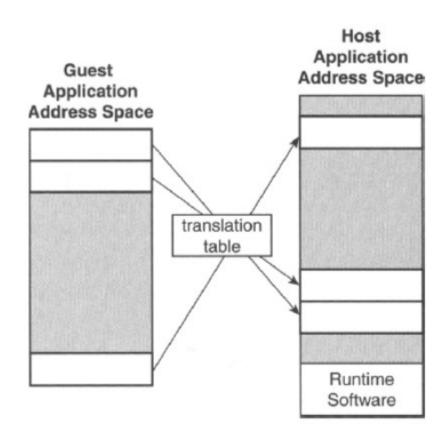
- System (full) virtualization
- Kernel emulation (wine/Windows, linux).

While many of current implementations are likely to be using the first flavour, we'll focus on the later one, which promises great speed enhancements, and a much greater attack surface...;)

Qemu architecture



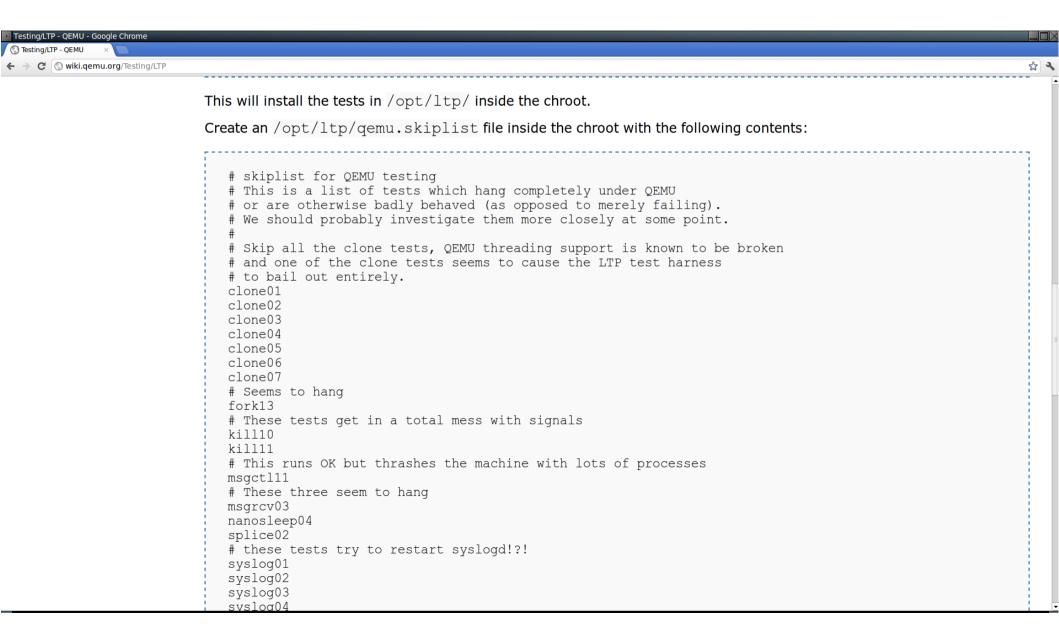
Understanding binary translation & memory sharing



Party time !!



Qemu regression tests...



Demos



Conclusion

Interresting technologies. Cool hacking tools, usefull for researchers.

There is room for massive security problems, the devil being in the details. Imho, not ready for Enterprise grade deployment.

No such a thing as third party assessment afaict.

From a strict game theory pov, your best interrest is probably to have <u>others</u> use that, but stay away from such technologies...

The Katsuni-Grothendieck theorem applied to silver bullet 0day sandboxing:

« You can't have it both ways! »

Thanks for inviting me.

Questions?

