

Breaking virtualization by switching to Virtual 8086 mode

Jonathan Brossard
CTO - P1 Code Security

HACKITO ERGO SUM



jonathan@p1sec.com
endrazine@gmail.com

Agenda

Virtualization : big picture

Attack surface analysis

Introducing the Virtual 8086 mode

Practical use : Fuzzing using vm86()

Virtualization : big picture

Market shares
Definitions

Virtualization : market shares

Source : Forrester Research 2009

**78% of companies have production
servers virtualized.**

20% only have virtualized servers.

Virtualization : market shares

Source : Forrester Research 2009

**VMWare is present in 98% of the
companies.**

**Microsoft virtualization products are
used by 17%.**

Citrix/Xen is used by 10%.

Virtualization : Definitions

Virtualization

Virtualization is the name given to the simulation with higher level components, of lower level components.

NOTE: Virtualization of applications (as opposed to full Oses) is out of topic.

Virtualization : Definitions

Virtual Machine

A virtual machine (VM) is : "an efficient, isolated duplicate of a real machine".

-- Gerald J. Popek and Robert P. Goldberg (1974). "Formal Requirements for Virtualizable Third Generation Architectures", Communications of the ACM.

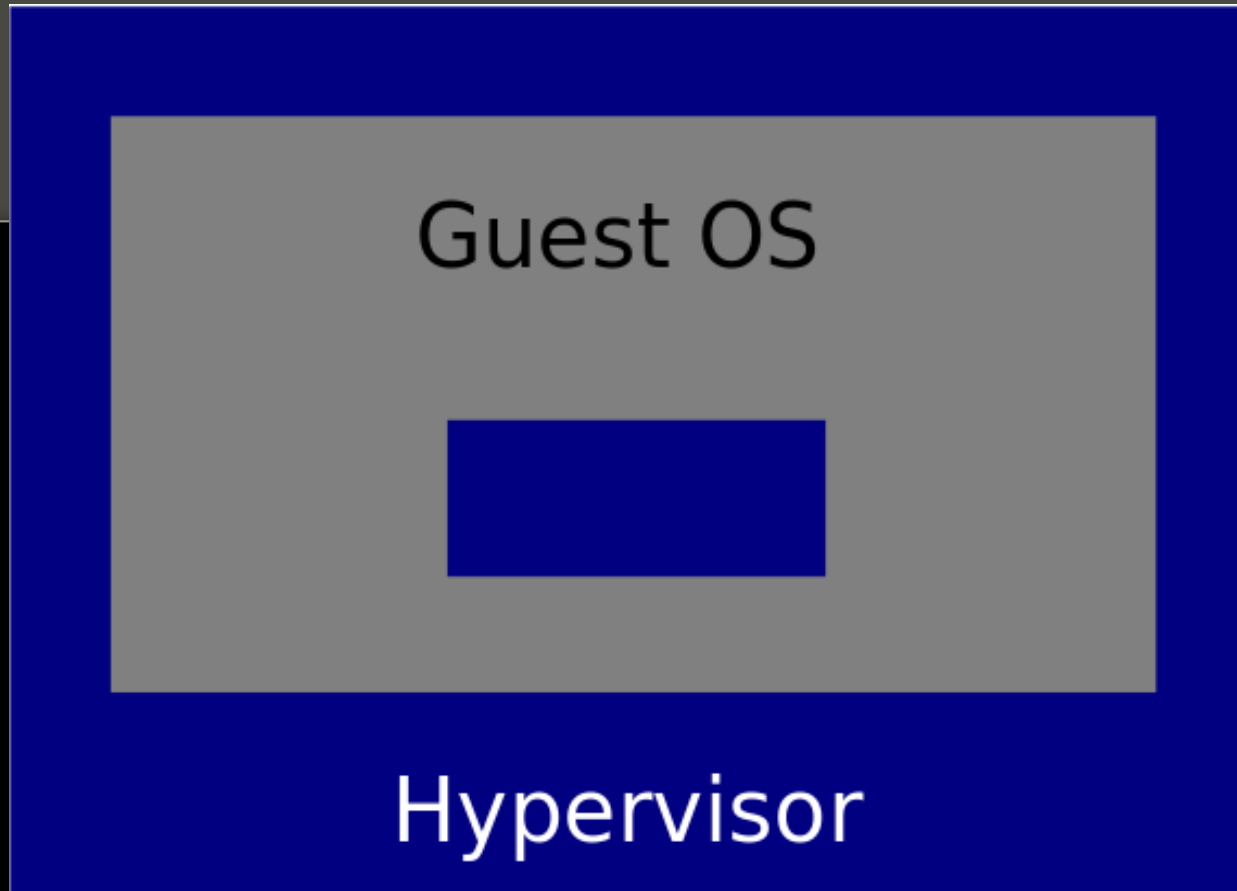
Virtualization : Definitions

Paravirtualization

**Requires the modification of the guest
Os (eg: Xen, UML, Qemu with
kqemu, VMWare Workstation with
VMWare Tools).**

Opposed to « full virtualization ».

Paravirtualization



Virtualization : Definitions

**There are two types of virtualizations :
Virtual Machine Monitors (or
Hypervisors) of type I and type II.**

Virtualization : Definitions

Hypervisors of type I

Run on bare metal (eg: Xen, Hyper-V, VMWare ESX).

Type I Hypervisor



The diagram illustrates the architecture of a Type I Hypervisor. It consists of three main layers: a top grey layer for the Guest OS, a middle blue layer for the Type I Hypervisor, and a bottom red layer for the Hardware. The Guest OS is contained within a grey rectangle, which is itself within a blue rectangle representing the hypervisor. The hypervisor layer is labeled 'Type I Hypervisor' in white text. The hardware layer is labeled 'Hardware' in black text.

Guest OS

Type I Hypervisor

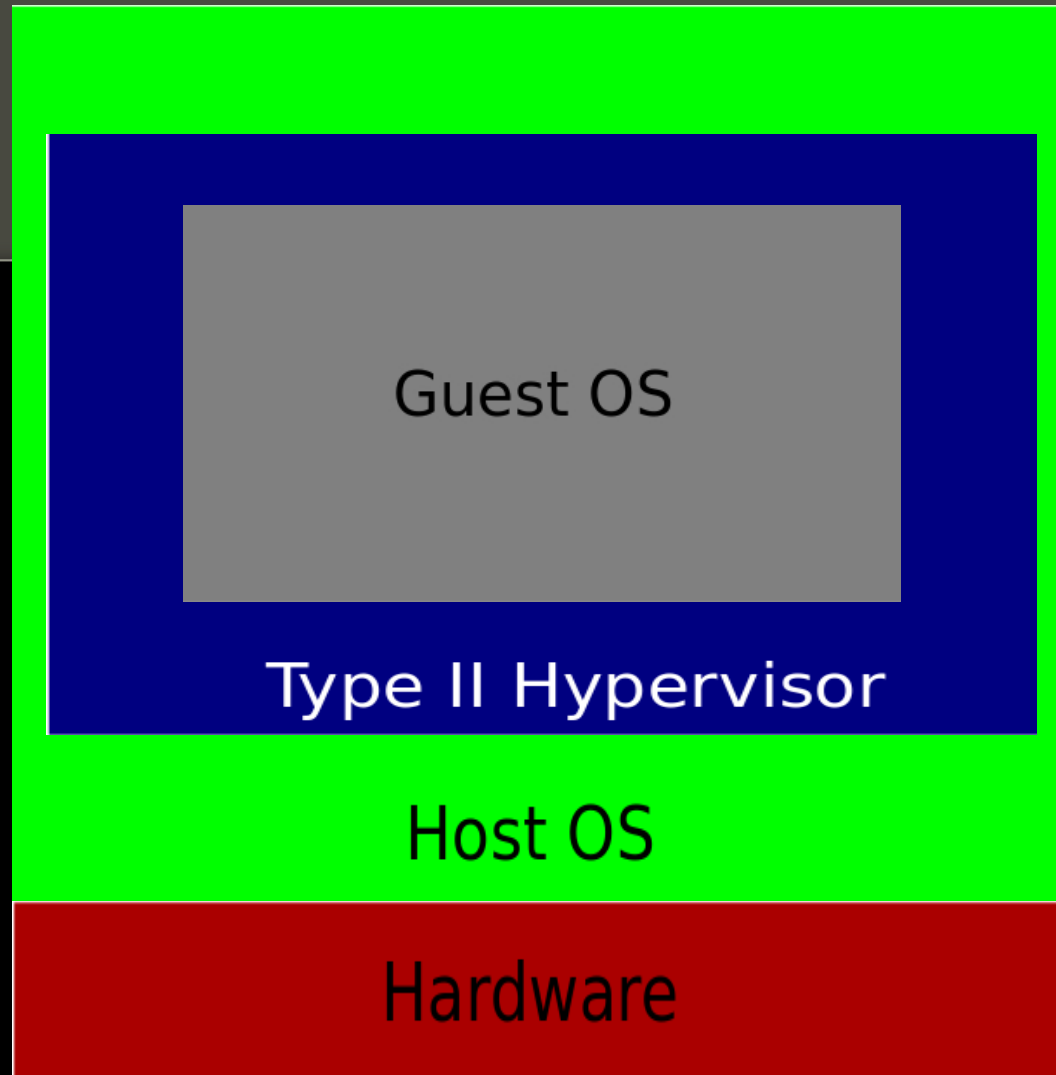
Hardware

Virtualization : Definitions

Hypervisors of type II

Run as a process inside a host OS to virtualize guests Oses (eg: Qemu, Virtualbox, VMWare Workstation, Parallels).

Type II hypervisor



Virtualization : Definitions

Isolation

Isolation of the userland part of the OS to simulate independant machines (eg: Linux-Vservers, Solaris « Zones », BSD « jails », OpenVZ under GNU/Linux).


Isolation

Userland
1

Userland
3

Kernel

Userland
2



Attack surface analysis

Privilege escalation on the host

VMware Tools HGFS Local Privilege Escalation Vulnerability

(<http://labs.iddefense.com/intelligence/vulnerabilities/display.php?id=712>)

Privilege escalation on the Guest

CVE-2009-2267 « Mishandled exception on page fault in VMware » Tavis Ormandy and Julien Tinnes

Attacking other guests

Vmare workstation guest isolation weaknesses (clipboard transfer)

<http://www.securiteam.com/securitynews/5GP021FKKO.html>

DoS (Host + Guests)

**CVE-2007-4591 CVE-2007-4593 (bad
ioctl's crashing the Host+Guests)**

Escape to host

**Rafal Wojtczuk (Invisible things,
BHUS 2008)**

**IDEFENSE VMware Workstation
Shared Folders Directory
Traversal Vulnerability
(CVE-2007-1744)**

(hardware level) attack vectors

Ioports:

**outb, outw, outl, outsb, outsw, outsl,
inb, inw, inl, insb, insw, insl, outb_p,
outw_p, outl_p, inb_p, inw_p, inl_p**

Problems: sequence, multiple ports

Ioctls:

int ioctl(int d, int request, ...)

Problems : arbitrary input size !

Introducing the Virtual 8086 mode

Introduced with Intel 386 (1985)

Introducing the Virtual 8086 mode

Intel x86 cpus support 3 modes

- Protected mode**
- Real mode**
- System Management Mode (SMM)**

Introducing the Virtual 8086 mode

Protected mode

This mode is the native state of the processor. Among the capabilities of protected mode is the ability to directly execute “real-address mode” 8086 software in a protected, multi-tasking environment. This feature is called virtual-8086 mode, although it is not actually a processor mode. Virtual-8086 mode is actually a protected mode attribute that can be enabled for any task.

Introducing the Virtual 8086 mode

Real-address mode

This mode implements the programming environment of the Intel 8086 processor with extensions (such as the ability to switch to protected or system management mode). The processor is placed in real-address mode following power-up or a reset.

Introducing the Virtual 8086 mode

System management mode (SMM)

This mode provides an operating system or executive with a transparent mechanism for implementing platform specific functions such as power management and system security. The processor enters SMM when the external SMM interrupt pin (SMI#) is activated or an SMI is received from the advanced programmable interrupt controller (APIC).

Nice things about Real mode / Virtual 8086 mode

Direct access to hardware via
interruptions !

example:

```
Mov ah, 0x42 ; read sector from drive
Mov ch, 0x01 ; Track
Mov cl, 0x02 ; Sector
Mov dh, 0x03 ; Head
Mov dl, 0x80 ; Drive (here first HD)
Mov bx, offset buff ; es:bx is destination

Int 0x13 ; hard disk operation
```

Complexity

$ax*bx*cx*dx$ (per interruption)

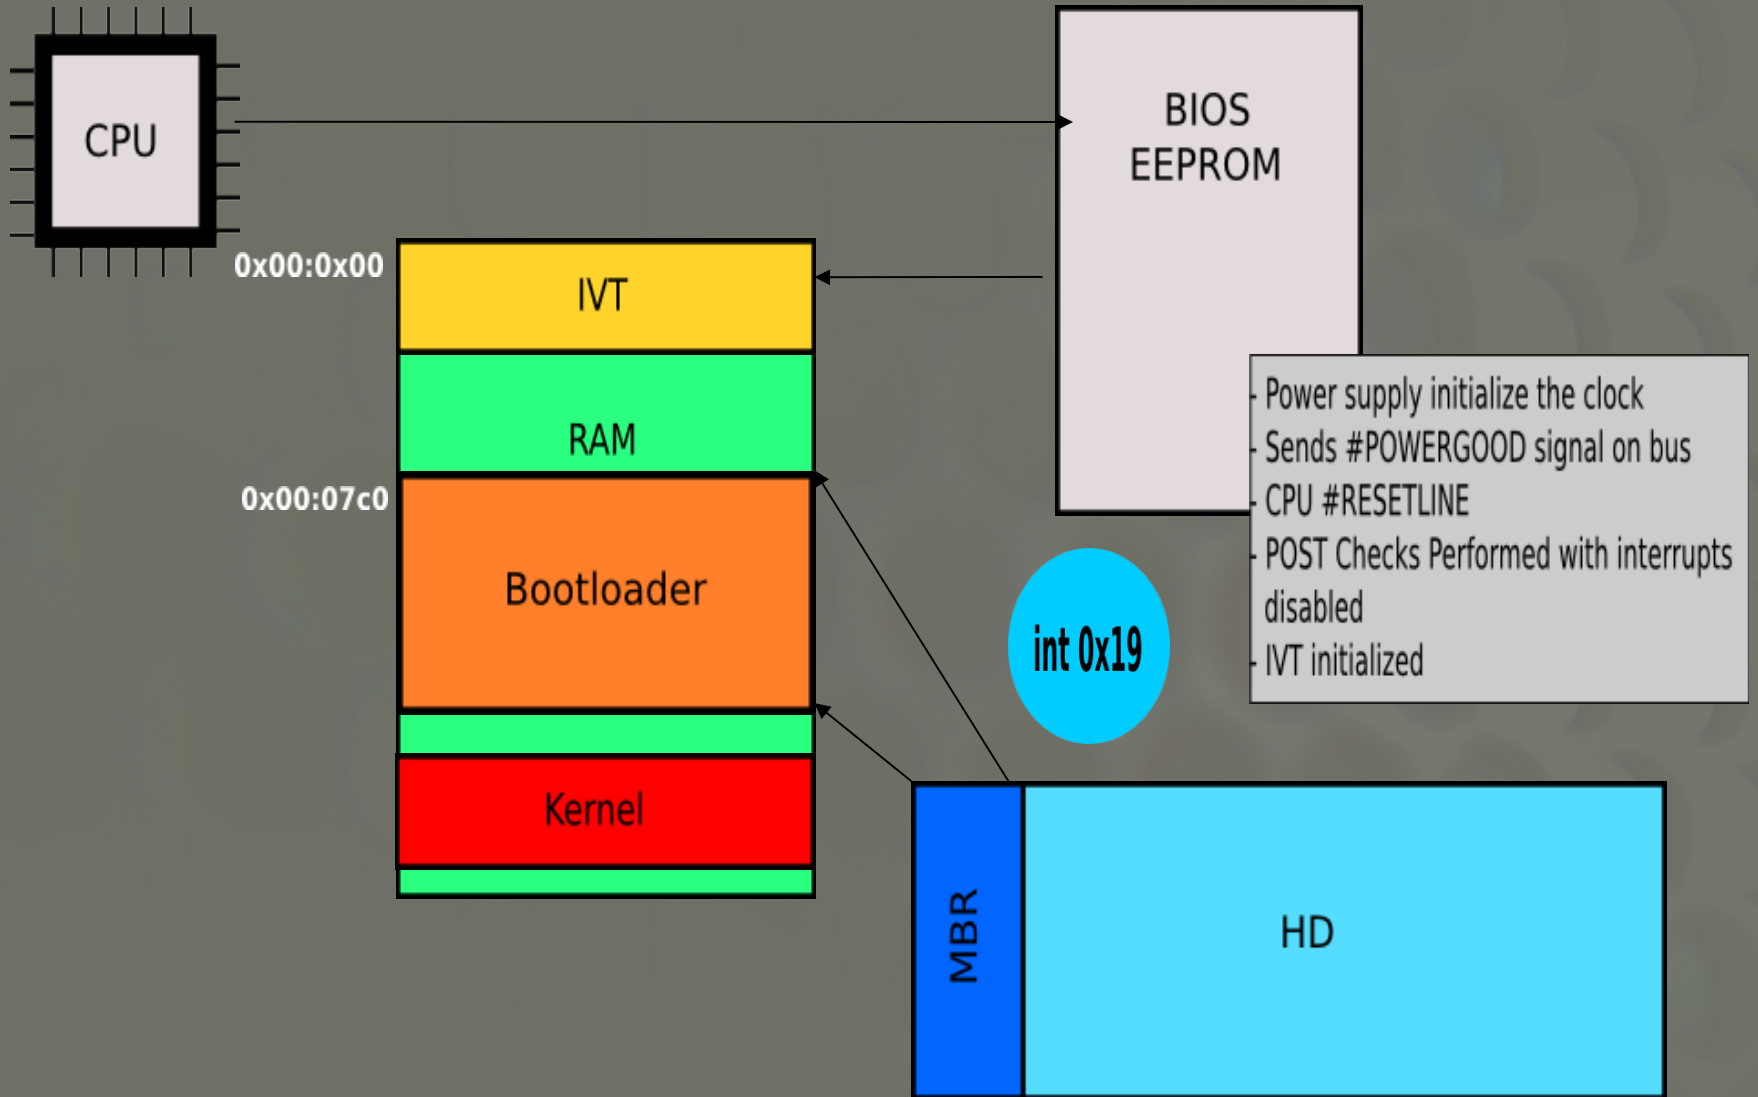
Id est: $[0;65535]^4 \sim 1.8 * 10^{19}$

=> still huge

=> much better than `ioctl()`'s arbitrary input length !

Introducing the Virtual 8086 mode

Putting it all together...



Introducing the Virtual 8086 mode

Corollary

The hypervisor runs under protected mode (ring0, ring1 (!!)) or ring3).

All of the guests run in protected mode.

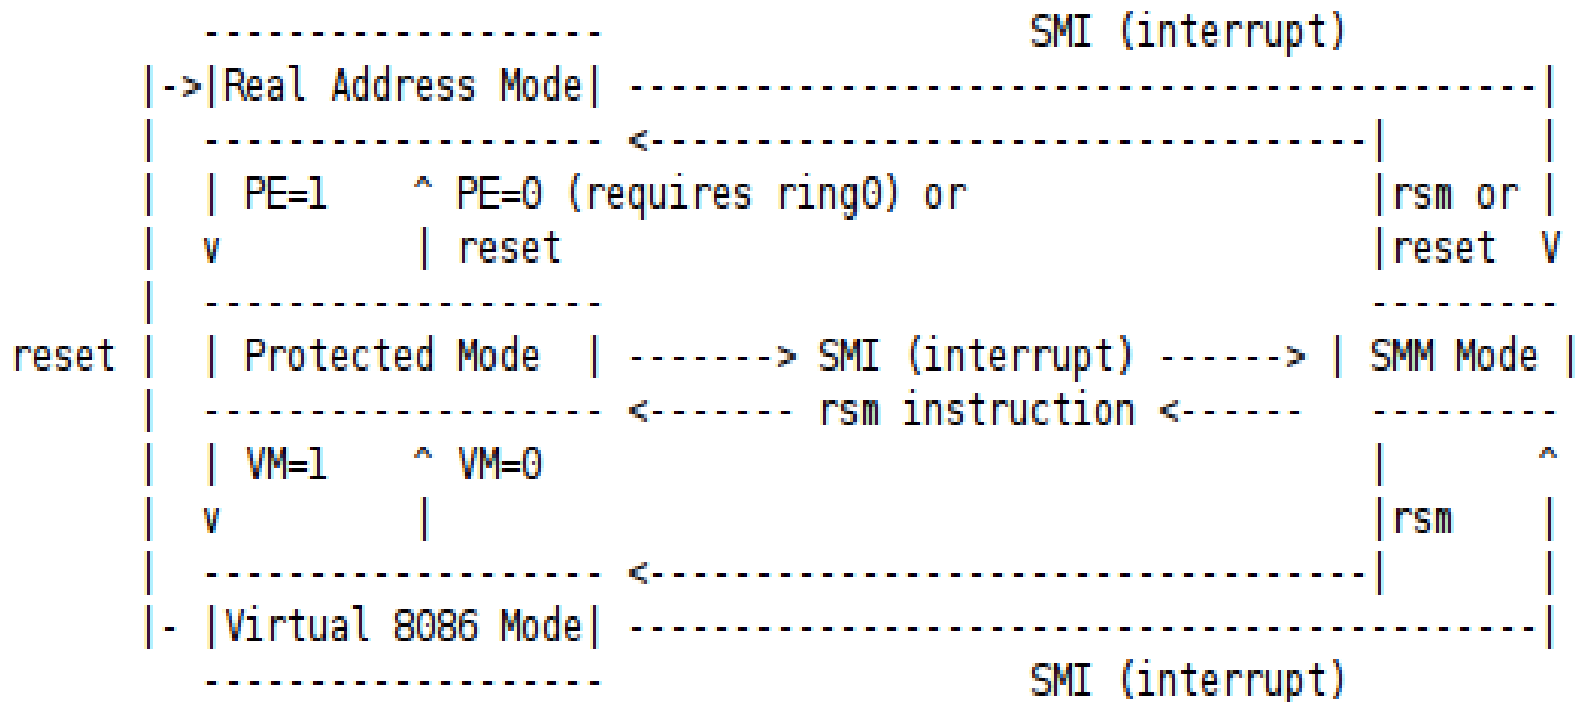
Introducing the Virtual 8086 mode

The kernel boots in (16b) real mode, and then switches to protected mode (32b).

The cpu normally doesn't get back to real mode until next reboot.

GAME OVER ?
Not quite ;)

Leaving protected mode ?



(Ascii Art : Courtesy of phrack 65)

Setting the VM flag in CR0 under protected mode would get us to Virtual Mode

Removing the PE flag from CR0 would get us back to real mode

Leaving protected mode ?

linux-2.6.31/arch/x86/kernel/reboot.c:

```
static const unsigned char real_mode_switch [] =
```

```
{
    0x66, 0x0f, 0x20, 0xc0,          /* movl %cr0,%eax */
    0x66, 0x83, 0xe0, 0x11,          /* andl $0x00000011,%eax */
    0x66, 0x0d, 0x00, 0x00, 0x00, 0x60, /* orl $0x60000000,%eax */
    /*
    0x66, 0x0f, 0x22, 0xc0,          /* movl %eax,%cr0 */
    0x66, 0x0f, 0x22, 0xd8,          /* movl %eax,%cr3 */
    0x66, 0x0f, 0x20, 0xc3,          /* movl %cr0,%ebx */
    0x66, 0x81, 0xe3, 0x00, 0x00, 0x00, 0x60, /* andl
    $0x60000000,%ebx */
    0x74, 0x02,                      /* jz f */
    0x0f, 0x09,                      /* wbinvd */
    0x24, 0x10,                      /* f: andb $0x10,al */
    0x66, 0x0f, 0x22, 0xc0          /* movl %eax,%cr0 */
};
```

Trouble is...

This obviously won't work inside a virtual machine !

Because CR[1-4] registers are themselves emulated

Truth is : we don't need to switch back to real mode/virtual 8086 mode !

Most Operating systems offer a way to run 16b applications (eg: MS DOS) under protected mode by emulating a switch to Virtual 8086 Mode.

Notably Windows (x86) and Linux (x86).

The Windows case

NTVDM : ntvdm.exe
« Windows 16b Virtual Machine »



Corbeille



Breaking
virtualization by...

```
Administrateur : Invite de commandes - command.com
Microsoft Windows [version 6.0.6002]
Copyright (c) 2006 Microsoft Corporation. Tous droits réservés.

C:\Users\Administrateur>command.com
Microsoft(R) Windows DOS
(C)Copyright Microsoft Corp 1990-2001.

C:\USERS\ADMINI~1>
```

Démarrer

Administrateur : Invit...

FR



19:01

The Linux case

The linux kernel provides an emulation of real mode in the form of two syscalls:

```
#define __NR_vm86old    113  
#define __NR_vm86      166
```

The Linux case

```
#include <sys/vm86.h>
```

```
int vm86old(struct vm86_struct *info);
```

```
int vm86(unsigned long fn, struct vm86plus_struct *v86);
```

```
struct vm86_struct {  
    struct vm86_regs regs;  
    unsigned long flags;  
    unsigned long screen_bitmap;  
    unsigned long cpu_type;  
    struct revector_struct  
        int_revector;  
    struct revector_struct  
        int21_revector;  
};
```

```
struct vm86_struct {
```

```
    struct vm86_regs regs;
```

```
    unsigned long flags;
```

```
    unsigned long screen_bitmap;
```

```
    unsigned long cpu_type;
```

```
    struct revector_struct
```

```
        int_revector;
```

```
    struct revector_struct
```

```
        int21_revector;
```

```
};
```

The Linux case

linux-2.6.31/arch/x86/include/asm/vm86.h:

```
struct vm86_regs {  
    long ebx;  
    long ecx;  
    long edx;  
    long esi;  
    long edi;  
    long ebp;  
    long eax;  
    (...)  
    unsigned short es, __esh;  
    unsigned short ds, __dsh;  
    unsigned short fs, __fsh;  
    unsigned short gs, __gsh;  
};
```

In a nutshell

- The switch to Virtual mode is completely emulated by the kernel (this will work inside a VM)
- We can still program using old school interruptions (easy !)
- Those interruptions are delivered to the hardware (id est: either the emulated one, or the real one).

=> We just got a « bare metal (possibly virtualized) hardware interface »

Practical use : Fuzzing
using vm86()

Looking at the IVT

Practical use : Fuzzing
using vm86()

Hypervisors bugs !

Virtualbox

```
00:21:13 !!
00:21:13
00:21:13 PANIC)
00:21:13
00:21:13 0000c0000
EIP=
00:21:13 ols:
00:21:13
00:21:13 _ZL1
00:21:13 Ex
00:21:13 ]
00:21:13
00:21:13
00:21:13
00:21:13
00:21:13 0000000
00:21:13 .esi=
```

Virtualbox (take 2)

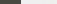
00:02:51 • Ubuntu Server [arrêté] - VirtualBox OSE

00:02:51 Machine Périphériques Aide

```
00:02:5Returned ecx: 134
```

00:02:51

VirtualBox - Guru Meditation

EIP=  Une erreur critique est survenue pendant l'exécution de la machine virtuelle et cette dernière a été suspendue.

00:02:51 🤝 a été suspendue.

00:02:51_ZL1 Pour trouver de l'aide allez à la section Community sur <http://www.virtualbox.org> ou voyez votre contrat de support. Veuillez fournir le fichier historique `VBox.log` et le fichier image `VBox.png` que vous trouverez dans le répertoire

00:02:5 /home/jonathan/.VirtualBox/Machines/Ubuntu Server/Logs ainsi qu'une description de
00:02:5 ce que vous faisiez quand l'erreur s'est produit. Vous pouvez également accéder aux fichiers
00:02:5 en sélectionnant **Afficher l'historique** dans le menu **Machine** de la fenêtre principale de
VirtualBox.

00:02:51 Activez le bouton **OK** si vous désirez arrêter la machine ou **Ignorer** pour la laisser telle quelle
00:02:51 pour le déboguage. Comme le déboguage nécessite des connaissances et des outils
00:02:51 spécialisés, il est conseillé de choisir **OK**.

00:02:51 OK Ignorer

.esi=

```
00:02:51 Returned eax: 129
```

```
Returned ebx: 157
```



T PANIC)

00000ab000

bols:


Ex
x]

00000000

rf nv up di

More bugs

Virtual PC (Guest ?)



Parallels (Guest)

----- Guest processor state -----

Inhibit Mask=0

CS=FF63 [0000FFFF 0000F30F] V=1

SS=FFD3 [0000FFFF 00CF9300] V=1

DS=0018 [0000FFFF 00CFF300] V=1

ES=0018 [0000FFFF 00CFF300] V=1

FS=FF9B [0000FFFF 00CF9300] V=1

GS=0018 [0000FFFF 00CF9300] V=1

EAX=000000A9 EBX=00005148 ECX=0000F686

EDX=0000000B

ESI=00002D72 EDI=000007E4 EBP=00002E99

ESP=00000FFA

EIP=0000FE96 EFLAGS=00023202

DEMOS

Thank you for coming

Questions ?



P1 Security
Priority One Security